

04_crnn_training

March 24, 2026

1 04 CRNN Training

This notebook walks through training the CRNN (Convolutional-Recurrent) OCR model on the extracted ground-truth line crops. The model uses a CNN backbone to extract visual features and a bidirectional LSTM to decode character sequences via CTC loss. Rare characters (diacritics, ligatures specific to Spanish renaissance print) receive higher weight in the loss via inverse-frequency weighting.

```
[7]: import sys
from pathlib import Path

cwd = Path.cwd().resolve()
ROOT = cwd if (cwd / 'src').exists() else cwd.parent
if str(ROOT) not in sys.path:
    sys.path.insert(0, str(ROOT))

import torch
import matplotlib.pyplot as plt

from src.crnn_model import build_vocab, compute_char_weights, CRNN
from src.utils import read_jsonl

GROUND_TRUTH_FILE = ROOT / 'data' / 'ground_truth' / 'ground_truth.jsonl'
CROP_DIR = ROOT / 'data' / 'crops'
OUTPUT_DIR = ROOT / 'outputs' / 'crnn'
```

1.1 1. Build vocabulary and compute character weights

We first collect all characters that appear in the ground truth and compute their frequency. Rare characters (e.g. ñ, ü, period-specific abbreviations) get **higher** weight in the CTC loss, which forces the model to pay more attention to them during training.

```
[8]: records = read_jsonl(GROUND_TRUTH_FILE)
texts = [r['text'] for r in records]
vocab = build_vocab(texts)
weights = compute_char_weights(texts, vocab)

print(f'Vocabulary size: {len(vocab)} characters')
```

```

print(f'Sample entries: {dict(list(vocab.items())[:10])}')

# Show the top 15 highest-weighted (rarest) characters
sorted_weights = sorted(
    [(ch, weights[idx].item()) for ch, idx in vocab.items()],
    key=lambda x: -x[1]
)[:15]
print('\nRarest characters (highest CTC weight):')
for ch, w in sorted_weights:
    print(f' {repr(ch):12s} weight={w:.3f}')

```

Vocabulary size: 67 characters

Sample entries: {'<BLANK>': 0, '\n': 1, ' ': 2, '!': 3, '&': 4, '(': 5, ')': 6, ',': 7, '-': 8, '.': 9}

Rarest characters (highest CTC weight):


'!'	weight=4.226
'('	weight=4.226
')'	weight=4.226
'0'	weight=4.226
'6'	weight=4.226
'H'	weight=4.226
'Q'	weight=4.226
'Ñ'	weight=4.226
'&'	weight=2.817
'1'	weight=2.817
'é'	weight=2.817
'í'	weight=2.817
'J'	weight=2.113
'Y'	weight=2.113
'ó'	weight=2.113

```

[9]: # Visualise the full weight distribution
chars = [ch for ch, _ in sorted(vocab.items(), key=lambda kv: kv[1])]
weight_vals = [weights[vocab[ch]].item() for ch in chars]

plt.figure(figsize=(14, 4))
plt.bar(range(len(chars)), weight_vals)
plt.xticks(range(len(chars)), chars, fontsize=7, rotation=90)
plt.title('Per-character CTC loss weights (inverse frequency)')
plt.ylabel('Weight')
plt.xlabel('Character')
plt.tight_layout()
plt.show()

```



1.2 2. Inspect the CRNN architecture

The model maps an image of shape (B, 1, H, W) through: 1. **CNN backbone** – four conv-BN-LeakyReLU-MaxPool blocks that collapse the height to 1 2. **BiLSTM** – two layers treating width positions as time steps 3. **Linear projection** – maps to log-probabilities over vocabulary

```
[10]: model = CRNN(num_classes=len(vocab))
total_params = sum(p.numel() for p in model.parameters())
print(f'Total parameters: {total_params:,}')
print(model)

# Smoke test: forward pass on random input
x = torch.randn(2, 1, 32, 256)
out = model(x)
print(f'\nInput shape: {tuple(x.shape)}')
print(f'Output shape: {tuple(out.shape)} (T, B, num_classes)')
```

Total parameters: 4,739,075

```
CRNN(
  (cnn): CNNBackbone(
    (blocks): Sequential(
      (0): Sequential(
        (0): Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): LeakyReLU(negative_slope=0.2, inplace=True)
        (3): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
      )
      (1): Sequential(
        (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): LeakyReLU(negative_slope=0.2, inplace=True)
```

```

        (3): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
            ceil_mode=False)
    )
    (2): Sequential(
      (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (2): LeakyReLU(negative_slope=0.2, inplace=True)
    )
    (3): Sequential(
      (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (2): LeakyReLU(negative_slope=0.2, inplace=True)
    )
  )
  (height_pool): AdaptiveAvgPool2d(output_size=(1, None))
)
(lstm): LSTM(512, 256, num_layers=2, bidirectional=True)
(classifier): Linear(in_features=512, out_features=67, bias=True)
)

```

Input shape: (2, 1, 32, 256)
Output shape: (64, 2, 67) (T, B, num_classes)

1.3 3. Run training

The training script handles dataset loading, batching, and saving the best checkpoint. Run it from the project root in a terminal:

```

python scripts/train_crnn.py \
  --ground-truth-file data/ground_truth/ground_truth.jsonl \
  --line-image-dir data/crops \
  --output-dir outputs/crnn \
  --epochs 30

```

Or call it directly from this notebook (blocking):

```

[11]: import subprocess, sys
result = subprocess.run(
    [
        sys.executable, str(ROOT / 'scripts' / 'train_crnn.py'),
        '--ground-truth-file', str(GROUND_TRUTH_FILE),
        '--line-image-dir', str(CROP_DIR),
        '--output-dir', str(OUTPUT_DIR),
        '--epochs', '30',
    ],

```

```
capture_output=True, text=True,
)
print(result.stdout)
if result.returncode != 0:
    print('STDERR:', result.stderr)
```

Device: cpu

Vocabulary size: 67 characters → /Users/ray/CNN-RNN/outputs/crnn/crnn_vocab.json

Dataset size: 5 samples

Epoch	1/30	loss=0.0000
Epoch	2/30	loss=0.0000
Epoch	3/30	loss=0.0000
Epoch	4/30	loss=0.0000
Epoch	5/30	loss=0.0000
Epoch	6/30	loss=0.0000
Epoch	7/30	loss=0.0000
Epoch	8/30	loss=0.0000
Epoch	9/30	loss=0.0000
Epoch	10/30	loss=0.0000
Epoch	11/30	loss=0.0000
Epoch	12/30	loss=0.0000
Epoch	13/30	loss=0.0000
Epoch	14/30	loss=0.0000
Epoch	15/30	loss=0.0000
Epoch	16/30	loss=0.0000
Epoch	17/30	loss=0.0000
Epoch	18/30	loss=0.0000
Epoch	19/30	loss=0.0000
Epoch	20/30	loss=0.0000
Epoch	21/30	loss=0.0000
Epoch	22/30	loss=0.0000
Epoch	23/30	loss=0.0000
Epoch	24/30	loss=0.0000
Epoch	25/30	loss=0.0000
Epoch	26/30	loss=0.0000
Epoch	27/30	loss=0.0000
Epoch	28/30	loss=0.0000
Epoch	29/30	loss=0.0000
Epoch	30/30	loss=0.0000

Training complete. Best loss: 0.0000

Checkpoint saved to: /Users/ray/CNN-RNN/outputs/crnn/crnn_checkpoint.pt

Vocabulary saved to: /Users/ray/CNN-RNN/outputs/crnn/crnn_vocab.json

1.4 4. Inference with the trained CRNN

Once training is complete, load the checkpoint and run a quick inference check.

```
[12]: from src.crnncr_model import CRNNOCRModel
from src.utils import CRNNConfig
from PIL import Image

checkpoint_path = OUTPUT_DIR / 'crnn_checkpoint.pt'
vocab_path = OUTPUT_DIR / 'crnn_vocab.json'

if checkpoint_path.exists() and vocab_path.exists():
    crnn = CRNNOCRModel(CRNNConfig(
        checkpoint_path=str(checkpoint_path),
        vocab_path=str(vocab_path),
    ))
    # Run on first available crop
    crop_paths = sorted(CROP_DIR.glob('*_crop.png'))
    if crop_paths:
        img = Image.open(crop_paths[0]).convert('L')
        text = crnn.predict_page_text([img])
        print(f'Page: {crop_paths[0].name}')
        print('CRNN output:')
        print(text[:500])
    else:
        print('No checkpoint found. Run training first.')
```

Page: buendia_instruccion_page_0001_crop.png

CRNN output:

sj